

## Under Sampling

```
scala> :paste
// Entering paste mode (ctrl-D to finish)

import org.apache.spark.sql.functions._
import org.apache.spark.sql.types._
import org.apache.spark.sql.expressions.Window
import org.apache.spark.ml.feature.{VectorAssembler, StringIndexer}
import org.apache.spark.ml.Pipeline
import org.apache.spark.ml.classification.{RandomForestClassificationModel, RandomForestClassifier}
import org.apache.spark.ml.tuning.{CrossValidator, CrossValidatorModel, ParamGridBuilder}
import org.apache.spark.ml.evaluation.{MulticlassClassificationEvaluator}
import org.apache.spark.ml.param.ParamMap
import org.apache.spark.sql.types.{IntegerType, DoubleType}
import org.apache.spark.sql.DataFrame

val schema_covid19 = StructType( StructField(" id", IntegerType, nullable = true) ::
    StructField("Assigned ID", IntegerType, nullable = true) ::
    StructField("Outbreak Associated", StringType, nullable = true) ::
    StructField("Age Group", StringType, nullable = true) ::
    StructField("Neighbourhood Name", StringType, nullable = true) ::
    StructField("FSA", StringType, nullable = true) ::
    StructField("Source of Infection", StringType, nullable = true) ::
    StructField("Classification", StringType, nullable = true) ::
    StructField("Episode Date", DateType, nullable = true) ::
    StructField("Reported Date", DateType, nullable = true) ::
    StructField("Client Gender", StringType, nullable = true) ::
    StructField("Outcome", StringType, nullable = true) ::
    StructField("Currently Hospitalized", StringType, nullable = true) ::
    StructField("Currently in ICU", StringType, nullable = true) ::
    StructField("Currently Intubated", StringType, nullable = true) ::
    StructField("Ever Hospitalized", StringType, nullable = true) ::
    StructField("Ever in ICU", StringType, nullable = true) ::
    StructField("Ever Intubated", StringType, nullable = true) :: Nil )

val raw_covid19_df = spark.read.format("csv").
    option("header", value = true).option("delimiter", ",").option("mode", "DROPMALFORMED").
    schema(schema_covid19).load("hdfs://BigData/Covid19Cases.csv").cache()

raw_covid19_df.printSchema()
raw_covid19_df.show(2)
```

```
raw_covid19_df.printSchema()
raw_covid19_df.show(2)

for( i <- raw_covid19_df.columns){
    println(i+" "+raw_covid19_df.filter(raw_covid19_df(i).isNull || raw_covid19_df(i) == "").count()+ " "+
    100*(raw_covid19_df.filter(raw_covid19_df(i).isNull || raw_covid19_df(i) == "").count()/(raw_covid19_df.count()))
}

val covid19_df = raw_covid19_df.filter(col("Outcome").isin("RESOLVED","FATAL"))
    .filter(col("Age Group").isNotNull)

covid19_df.count()

val indexer = new StringIndexer()
    .setInputCol("Outcome")
    .setOutputCol("OutcomeIDX")

print(indexer)

val covid19_df1 = indexer.fit(covid19_df).transform(covid19_df)

/*import scala.collection.mutable.ListBuffer*/
var f_indexers = new Array[org.apache.spark.ml.PipelineStage](0)
val featuresList = List("Outbreak Associated","Age Group","Source of Infection","Client Gender","Ever Hospitalized",
    "Ever in ICU","Ever Intubated")

for (feature <- featuresList){
    print(feature)
    val f_indexer = new StringIndexer().setInputCol(feature).setOutputCol(feature+ " IDX")
    print(f_indexer)
    f_indexers = f_indexers :+ f_indexer
}

f_indexers

val fpipeline = new Pipeline()
    .setStages(f_indexers)

val covid19_df2= fpipeline.fit(covid19_df1).transform(covid19_df1)

covid19_df2.filter(col("Outcome")==="FATAL").show(2)
```

```
giranabhi@bigdata-m: ~ -- Mozilla Firefox
https://ssh.cloud.google.com/projects/iron-zodiac-347119/zones/us-central1-b/instances/bigdata-m?authuser=0&hl=en_US&projectNumber=782503674315&useAdminProxy=true&trou

val fpipeline = new Pipeline()
.setStages(f_indexers)

val covid19_df2= fpipeline.fit(covid19_df1).transform(covid19_df1)
covid19_df2.filter(col("Outcome")==="FATAL").show(2)

val resolved_df = covid19_df2.filter(col("Outcome") == "RESOLVED")
val fatal_df = covid19_df2.filter(col("Outcome") == "FATAL")
val ratio:Double = (resolved_df.count()/fatal_df.count())

// Under Sampling 'resolved' records and combining with 'fatal' population in order to generate a balanced dataset
val r:Double = (1/ratio).toDouble
print(r)
val sampled_resolved_df = resolved_df.sample(false,r, 42)
val combined_sampled_df = sampled_resolved_df.unionAll(fatal_df)
combined_sampled_df.filter(col("Outcome")==="FATAL").count()

val assembler = new VectorAssembler()
.setInputCols(Array("Outbreak Associated IDX","Age Group IDX","Source of Infection IDX","Client Gender IDX","Ever Hospitalized IDX",
                    "Ever in ICU IDX","Ever Intubated IDX"))
.setOutputCol("assembled-features")

val rf = new RandomForestClassifier()
.setFeaturesCol("assembled-features")
.setLabelCol("OutcomeIDX")
.setSeed(42)

val pipeline = new Pipeline()
.setStages(Array(assembler, rf))

val evaluator = new MulticlassClassificationEvaluator()
.setLabelCol("OutcomeIDX")
.setPredictionCol("prediction")
.setMetricName("accuracy")

val paramGrid = new ParamGridBuilder()
.addGrid(rf.maxDepth, Array(3, 4))
.addGrid(rf.impurity, Array("entropy","gini")).build()

val cross_validator = new CrossValidator()
```

```
giranabhi@bigdata-m: ~ -- Mozilla Firefox
https://ssh.cloud.google.com/projects/iron-zodiac-347119/zones/us-central1-b/instances/bigdata-m?authuser=0&hl=en_US&projectNumber=782503674315&useAdminProxy=true&trou

.addGrid(rf.impurity, Array("entropy","gini")).build()

val cross_validator = new CrossValidator()
.setEstimator(pipeline)
.setEvaluator(evaluator)
.setEstimatorParamMaps(paramGrid)
.setNumFolds(3)

/* val trainData = covid19_weighted.sample("OutcomeIDX", fractions=(0.0: 0.09, 1.0: 0.7), seed=42)*/
val Array(trainingData, testData) = combined_sampled_df.randomSplit(Array(0.8, 0.2), 40)

val cvModel = cross_validator.fit(trainingData)

val predictions = cvModel.transform(testData)

val accuracy = evaluator.evaluate(predictions)
print(accuracy)

// Exiting paste mode, now interpreting.

<console>:90: warning: a pure expression does nothing in statement position; multiline expressions may require enclosing parentheses
    f_indexers
    ^
root
|-- _id: integer (nullable = true)
|-- Assigned ID: integer (nullable = true)
|-- Outbreak Associated: string (nullable = true)
|-- Age Group: string (nullable = true)
|-- Neighbourhood Name: string (nullable = true)
|-- FSA: string (nullable = true)
|-- Source of Infection: string (nullable = true)
|-- Classification: string (nullable = true)
|-- Episode Date: date (nullable = true)
|-- Reported Date: date (nullable = true)
|-- Client Gender: string (nullable = true)
|-- Outcome: string (nullable = true)
|-- Currently Hospitalized: string (nullable = true)
|-- Currently in ICU: string (nullable = true)
|-- Currently Intubated: string (nullable = true)
|-- Ever Hospitalized: string (nullable = true)
|-- Ever in ICU: string (nullable = true)
```

Windows taskbar showing search bar, taskbar icons (Edge, Mail, File Explorer, Chrome, WhatsApp, Teams, Firefox, OneDrive, Word), system tray (9°C, 20:33, 13-04-2022, ENG, 26).

Windows taskbar showing search bar, taskbar icons (Edge, Mail, File Explorer, Chrome, WhatsApp, Teams, Firefox, Photos, OneDrive, Word), system tray (9°C, network, volume, ENG, 20:33, 13-04-2022, 26).

```

|_id|Assigned_ID|Outbreak Associated|_Age Group|Neighbourhood Name|FSA|Source of Infection|Classification|Episode Date|Reported Date|Client Gender|Outcome|
|_id|Assigned_ID|Outbreak Associated|_Age Group|Neighbourhood Name|FSA|Source of Infection|Classification|Episode Date|Reported Date|Client Gender|Outcome|
|_id|Assigned_ID|Outbreak Associated|_Age Group|Neighbourhood Name|FSA|Source of Infection|Classification|Episode Date|Reported Date|Client Gender|Outcome|

```

```

|_id|Assigned_ID|Outbreak Associated|_Age Group|Neighbourhood Name|FSA|Source of Infection|Classification|Episode Date|Reported Date|Client Gender|Outcome|
|_id|Assigned_ID|Outbreak Associated|_Age Group|Neighbourhood Name|FSA|Source of Infection|Classification|Episode Date|Reported Date|Client Gender|Outcome|
|_id|Assigned_ID|Outbreak Associated|_Age Group|Neighbourhood Name|FSA|Source of Infection|Classification|Episode Date|Reported Date|Client Gender|Outcome|

```

only showing top 2 rows

```

0.9430455635491607import org.apache.spark.sql.functions._
import org.apache.spark.sql.types._
import org.apache.spark.sql.expressions.Window
import org.apache.spark.ml.feature.{VectorAssembler, StringIndexer}
import org.apache.spark.ml.Pipeline
import org.apache.spark.ml.classification.{RandomForestClassificationModel, RandomForestClassifier}
import org.apache.spark.ml.tuning.{CrossValidator, CrossValidatorModel, ParamGridBuilder}
import org.apache.spark.ml.evaluation.MulticlassClassificationEvaluator
import org.apache.spark.ml.param.ParamMap
import org.apache.spark.sql.types.{IntegerType, DoubleType}
import org.apache.spark.sql.DataFrame
import org.apache.spark.sql.types.StructType
schema_covid19: org.apache.spark.sql.types.StructType = StructType(StructField(_id,IntegerType,true), StructField(Assigned_ID,IntegerType,true...
scala> Connected, host fingerprint: ssh-rsa 0 77:91:A2:9F:E9:AB:18:62:E3:CA:A6:6B:7B:25:15:EB:A9:01:BD:F8:E8:B9:9E:AD:0D:40:E4:E8:FA:D0:5A:3D
Linux bigdata-m 5.10.0-0.bpo.12-amd64 #1 SMP Debian 5.10.103-1-bpo10+1 (2022-03-08) x86_64

```

The programs included with the Debian GNU/Linux system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/\*/\*copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent  
permitted by applicable law.  
Last login: Thu Apr 14 00:20:20 2022 from 35.235.245.130

giranabhi@bigdata-m:~\$

Type here to search



9°C

20:34

13-04-2022

ENG

26